

Comparing Classifiers (../comparing-classifiers/)

Contents

- Tutorial example
- The hardware
- MNIST
 - Neural Networks
 - SVM
 - Random Forest
 - k nearest neighbors
 - Decision Tree
 - Adaboost
 - Gradient Boosting
 - Naive Bayes
 - LDA
 - QDA
- MNIST Summary
- IRIS summary
- TL;DR

Classification problems occur quite often and many different classification algorithms have been described and implemented. But what is the best algorithm for a given error function and dataset?

I read questions like "I have problem X. What is the best classifier?" quite often and my first impulse is always to write: Just try them!

I guess people asking this question might think that it is super difficult to do so. However, the sklearn tutorial contains a very nice example where many classifiers are compared (source (http://scikit-learn.org/stable/auto_examples/classification/plot_classifier_comparison.html)).

This article gives you an overview over some classifiers:

- SVM (<http://scikit-learn.org/stable/modules/generated/sklearn.svm.SVC.html>)
- k-nearest neighbors (<http://scikit-learn.org/stable/modules/generated/sklearn.neighbors.KNeighborsClassifier.html>)
- Random Forest (<http://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestClassifier.html>)
- AdaBoost Classifier (<http://scikit-learn.org/stable/modules/generated/sklearn.ensemble.AdaBoostClassifier.html>)
- Gradient Boosting (<http://scikit-learn.org/stable/modules/generated/sklearn.ensemble.GradientBoostingClassifier.html>)
- Naive Bayes (http://scikit-learn.org/stable/modules/generated/sklearn.naive_bayes.GaussianNB.html)
- LDA (http://scikit-learn.org/0.16/modules/generated/sklearn.linear_model.LDA.html)
- QDA (http://scikit-learn.org/0.16/modules/generated/sklearn.linear_model.QDA.html)
- RBMs (http://scikit-learn.org/stable/modules/generated/sklearn.neural_network.BernoulliRBM.html)
- Logistic Regression (http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html)
- RBM (http://scikit-learn.org/stable/modules/generated/sklearn.neural_network.BernoulliRBM.html) + Logistic Regression Classifier

Of course, neural networks are also one very powerful ML classifier I may not forget. As sklearn does not have neural networks, I've installed `skflow` (<https://github.com/tensorflow/skflow>).

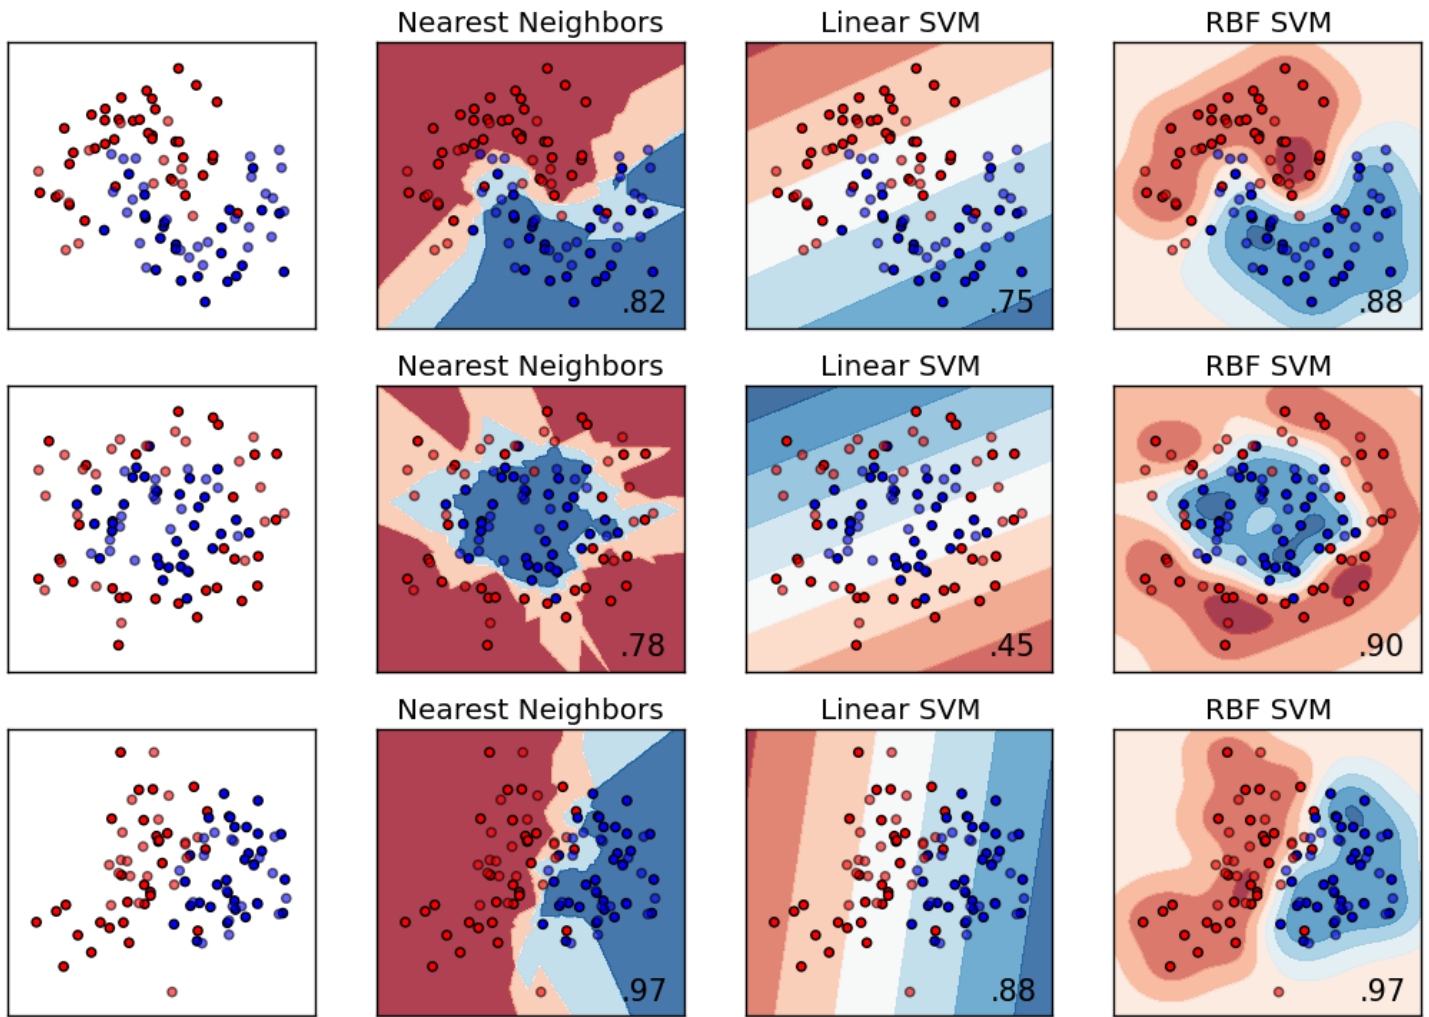
Tutorial example ¶

The sklearn tutorial creates three datasets with 100 points per dataset and 2 dimensions per point:

1. **Moons:** Two interleaving half-circles
2. **Circles:** A larger circle containing the smaller one
3. **Linear:** A linearly separable dataset

Each of those three datasets has added noise. This means for some points there might be no way of classifying them correctly.

Here are the results



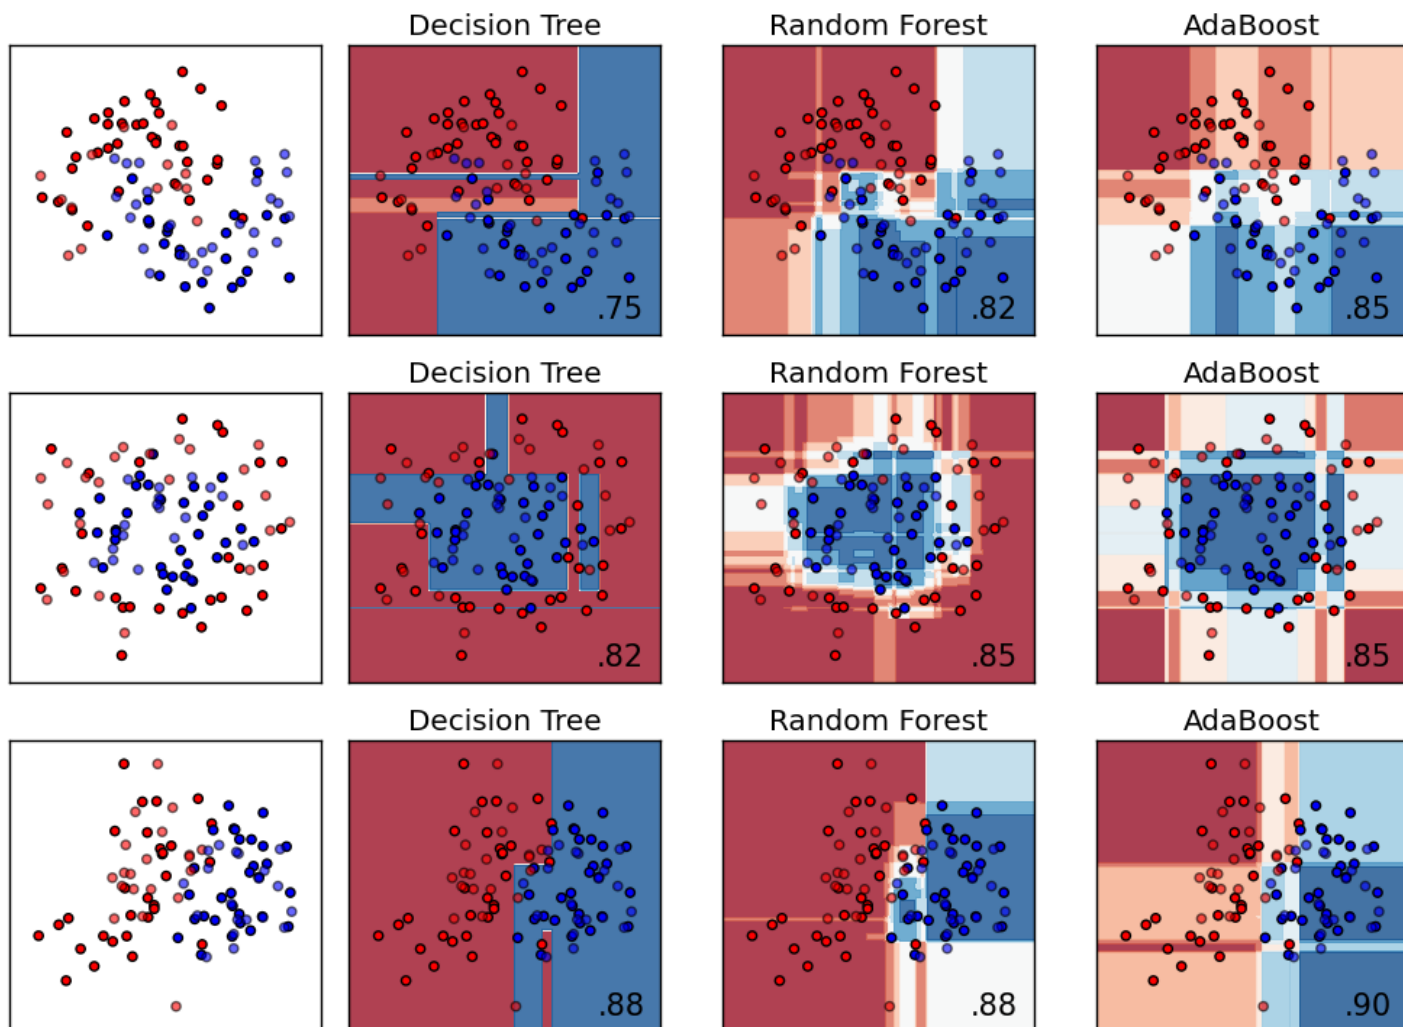
(../images/2016/01/ml-classifiers-1.png)

k nearest neighbors, linear and RBFSVM

One can see that k nearest neighbors gives arbitrary decision boundaries. Overall, they look reasonable. However, there are often strange zig-zag patterns.

The linear SVM in contrast has a very easy decision boundary: a line. It is no surprise that it can't deal with the moons dataset. Note that a random guess would be right in 50% of the cases.

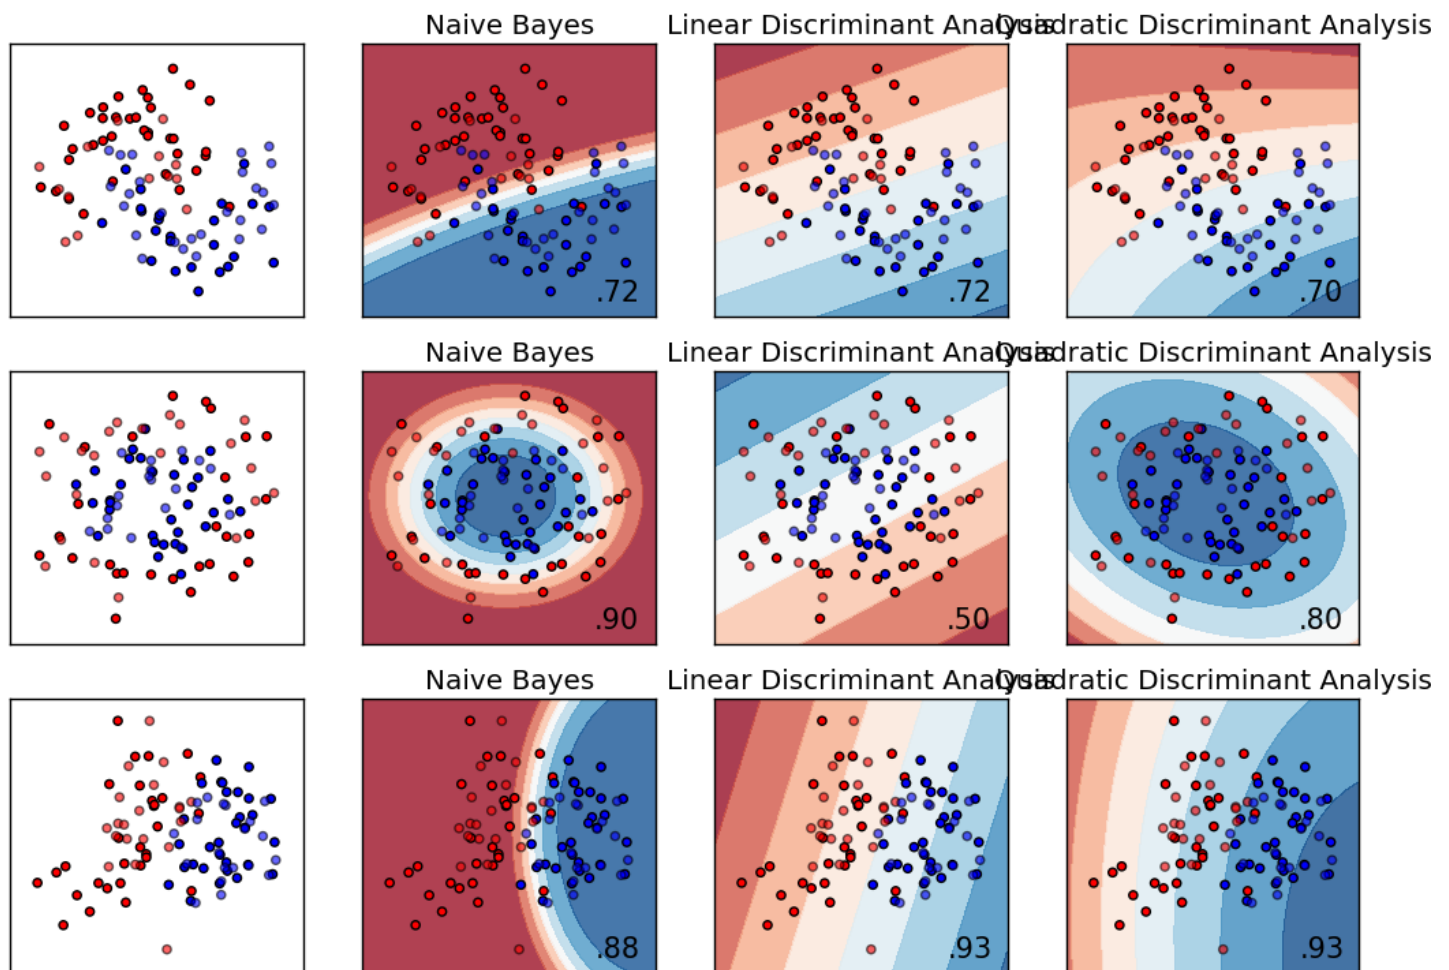
The RBF SVM has very nice decision boundary. It is smooth, matches the pattern and is able to adjust to all three examples.



(../images/2016/01/ml-classifiers-2.png)

Decision Tree, Random Forest, AdaBoost

Decision Trees, Decision Forests and AdaBoost all show very similar patterns. The boundaries change in parallel to the coordinate axes which looks very unnatural.



(../images/2016/01/ml-classifiers-3.png)

Naive Bayes, LDA, QDA

Naive Bayes shows nice, smooth patterns. However, those patterns seem to be a bit too simple. LDA is again linear (see linear SVM). Comparing QDA to Naive Bayes is interesting. Although they get similar performance for the first dataset, I would argue that the naive bayes classifier is much better as it is much more confident for its classification. Even more extreme is the last example. I'm astonished that the QDA gets 93% with that boundary; Naive Bayes seems to find a much better boundary.

The hardware ¶

The following comparison is done on a PC with an Intel i7-4820K CPU (http://ark.intel.com/de/products/77781/Intel-Core-i7-4820K-Processor-10M-Cache-up-to-3_90-GHz) and a NVIDIA GeForce GTX Titan Black GPU.

MNIST ¶

MNIST is a dataset of $28\text{px} \times 28\text{px}$ greyscale images. Each of the images contains a digit (0, 1, 2, 3, 4, 5, 6, 7, 8, 9). The task is to classify the image into one of the 10 digit classes.

Guessing randomly will give an accuracy of $\frac{1}{10} = 0.1$.

Neural Networks ¶

Please note that there are neural networks which get much better accuracy. Most notably the MNIST Expert tutorial (<https://www.tensorflow.org/versions/master/tutorials/mnist/pros/index.html#deep-mnist-for-experts>) with 99.2% accuracy.

Simple Network ¶

```

Classifier: NN 500:200
Training time: 79.5696s
Testing time: 0.3480s
Confusion matrix:
[[2248   1   5   1   2   4   8   2   5   2]
 [   1 2565  10   1   1   0   2   7   1   0]
 [   7   2 2258  14   5   0   9   6  10   3]
 [   0   0  12 2294   0  23   0   6   3  10]
 [   0   3   3   0 2161   0   8   5   1  30]
 [   4   1   1  16   1 2014  17   1   5   9]
 [  11   7   1   0   5   6 2237   0   4   0]
 [   3   6  14   7   3   1   0 2355  10  18]
 [   3   7   3  14   2  17   4   1 2161   3]
 [   4   4   0   4  16   8   0   7   6 2340]]
Accuracy: 0.9798

```

Dropout Network ¶

```

Classifier: NN 500:200 dropout
Training time: 118.2654s
Testing time: 0.3918s
Confusion matrix:
[[2250   1   7   1   1   1   5   4   4   4]
 [   1 2567   9   1   1   0   0   3   5   1]
 [   6   6 2272   3   2   1   3  10   8   3]
 [   0   0  26 2260   0  24   0  10  19   9]
 [   0   3   5   0 2152   0   7   3   1  40]
 [   8   3   3  12   2 1983  20   6  21  11]
 [  11   6   3   0   7   1 2237   0   6   0]
 [   2   7  13   3  11   0   1 2363   5  12]
 [   7   7   9   5   3   3   1   2 2170   8]
 [   3   3   1   3  13   2   0  19   8 2337]]
Accuracy: 0.9780

```

CNN ¶

```

Classifier: CNN
Training time: 391.8810s
Testing time: 1.2035s
Confusion matrix:
[[2243   0   5   0   0   5   9   1  12   3]
 [   1 2548  20   4   2   0   1   6   6   0]
 [   3   8 2253   9   3   1   4  17  14   2]
 [   0   4  13 2290   0  12   0   9  11   9]
 [   2   4   5   0 2164   0   8   5   3  20]
 [   6   2   3  15   0 2016   9   3   9   6]
 [  12  12   1   1   6   6 2227   0   6   0]
 [   3   4  11   3   4   1   0 2374   4  13]
 [   3  15   6  13   4  11   3   8 2145   7]
 [   6   5   0  11  16   8   0  24  13 2306]]
Accuracy: 0.9769

```

SVM ¶

There is a ton of literature / papers about SVMs (Support Vector Machines). I've summed up the basics on Using SVMs with sklearn (<https://martin-thoma.com/svm-with-sklearn/>).

I've trained two SVMs: A simple, linear one and one with an RBF kernel as I found it online (I'm sorry, I don't remember where I found those parameters :-/).

Please note the the SVM implementation of sklearn does not use the GPU. However, there are GPU implementations of SVMs (<http://fastml.com/running-things-on-a-gpu/>) around.

Linear SVM ¶

```

Classifier: linear SVM
Training time: 168.6950s
Testing time: 158.0101s
Confusion matrix:
[[2226    0    9    2    6   12    8    3   11    1]
 [   1 2537   18    3    3    1    1    7   17    0]
 [   12   16 2158   25   24    6   27   19   25    2]
 [    3    7   46 2188    4   47    3   18   27    5]
 [    2    5   19    1 2117    1    8    6    3   49]
 [   18   13   11   73   20 1872   31    0   26    5]
 [   20    6   22    1   10   30 2179    0    3    0]
 [    5   10   32   11   30    5    0 2268    5   51]
 [   11   39   26   47   10   40    7    7 2018   10]
 [   11    9    9   24   64    8    0   61   14 2189]]
Accuracy: 0.9416

```

Adjusted SVM ¶

```

Classifier: adj. SVM
Training time: 347.1539s
Testing time: 234.5724s
Confusion matrix:
[[2258    1    4    1    2    2    3    1    4    2]
 [    1 2566    9    1    1    0    0    7    3    0]
 [    4    1 2280    5    4    0    1    9    8    2]
 [    0    0   14 2304    1   13    0    6    8    2]
 [    2    2    2    0 2183    0    7    5    0   10]
 [    4    0    0   16    3 2026   12    1    4    3]
 [    7    5    3    0    5    2 2245    0    4    0]
 [    1    6   11    2    5    1    0 2373    5   13]
 [    3    9    4    9    4   10    2    3 2166    5]
 [    3    2    2    6   19    6    0   12   10 2329]]
Accuracy: 0.9840

```

Random Forest ¶

Data:

- `n_estimators=50`
- `n_jobs=10`

```

Classifier: Random Forest
Training time: 2.1359s
Testing time: 26.0763s
Confusion matrix:
[[2246    1    4    1    4    2    7    2   11    0]
 [    1 2543   18    5    5    2    3    7    4    0]
 [    7    2 2233   20    9    2    9   16   14    2]
 [    0    3   36 2240    0   20    3   16   19   11]
 [    3    1    5    0 2142    1   11    3    7   38]
 [    7    4    4   30    6 1977   16    3   14    8]
 [   13   11    4    0   10   15 2210    0    8    0]
 [    3    8   29    2   19    0    0 2315    7   34]
 [    3   12   18   17    9   26    4    7 2103   16]
 [   10    6    6   24   27   13    3   20   18 2262]]
Accuracy: 0.9641

```

Alternatively:

- `max_depth=5`
- `n_estimators=10`
- `max_features=1`

```

Classifier: Random Forest 2
Training time: 0.2077s
Testing time: 22.2770s
Confusion matrix:
[[1955   32   63   64   12    4  109   21   13    5]
 [    1 2524   20   14    1    6   10    6    6    0]
 [  252  425 1198  151   64    1  145   15   55    8]
 [  136  195  140 1641   28   11   22   95   65   15]
 [   92  320   21   45 1199    9   76  153    8  288]
 [  312  383   67  655   78  268   47   94  134   31]
 [  199  364  125   58   96   13 1408    5    2    1]
 [   83  424   10   70  101    1   19 1555   56   98]
 [  392  574   44  147   52   17   71  106  773   39]
 [   71  338   11   43  579    2    8  632   24  681]]
Accuracy: 0.5715

```

k nearest neighbors ¶

```

Classifier: k nn
Training time: 4.6439s
Testing time: 1261.7815s
Confusion matrix:
[[2260    1    4    0    0    1    6    2    2    2]
 [    0 2572    5    0    0    0    1    8    1    1]
 [   16   15 2235    9    1    0    5   26    5    2]
 [    2    5   14 2276    0   27    1    8    9    6]
 [    4   19    0    0 2131    0    8    4    0  45]
 [   10    5    3   28    5 1977   25    2    4   10]
 [   12    9    0    0    4    7 2239    0    0    0]
 [    1   18    4    1   12    0    0 2349    3   29]
 [   11   32    8   36   11   34    5    7 2053   18]
 [    6    8    4   14   26    4    0   19    5 2303]]
Accuracy: 0.9695

```

Decision Tree ¶

Data:

- max_depth=5

```

Classifier: Decision Tree
Training time: 3.1346s
Testing time: 0.0313s
Confusion matrix:
[[1767    0   11   25   12  120  137   71  114   21]
 [    1 2065  128  108   13   17   41   66  131   18]
 [   42   44 1248   37  121   21  227   76  339  159]
 [   33   22   32 1484   33  107   52   81  266  238]
 [    0   15   45   33 1284   42   42   45  213  492]
 [   42   10   21  229  166  577  137  123  254  510]
 [   34   33   66   24  103   65 1734   24  102   86]
 [   10   14  179   57   53   21   19 1775   79  210]
 [    1   98  129   43   43   42  160   29 1439  231]
 [    4    8   86   59  125   95   36   75  167 1734]]
Accuracy: 0.6540

```

Adaboost ¶

You should note that you can use arbitrary base classifiers with Adaboost. The default ones of `sklearn.ensemble.AdaBoostClassifier` (<http://scikit-learn.org/stable/modules/generated/sklearn.ensemble.AdaBoostClassifier.html>) is `sklearn.tree.DecisionTreeClassifier` (<http://scikit-learn.org/stable/modules/generated/sklearn.tree.DecisionTreeClassifier.html>)

```

Classifier: AdaBoost
Training time: 37.6443s
Testing time: 1.5815s
Confusion matrix:
[[1994    0    75     8     6   113    51     3    15    13]
 [    0 2435    27    22     2    10    12    37    42     1]
 [   97   39 1341    85    39    38   416    39   196    24]
 [  108   52   37 1508    13   313    66    64   122    65]
 [   11   16   48   23 1662    49   23   134    90   155]
 [   81   56   30   309    51 1255    57    17   129    84]
 [   29   28   151     7    80   43 1914     2    17     0]
 [   25   37   33   36   70   30     0 1761    37   388]
 [   30   80   48   215    16   85   30    19 1615    77]
 [   13   29   68   66   356    74     1   171    78 1533]]
Accuracy: 0.7367

```

Gradient Boosting ¶

Gradient boosting with `xgboost` has won in the Rossmann Store Sales prediction (source (<http://blog.kaggle.com/2015/12/21/rossmann-store-sales-winners-interview-1st-place-gert/>)).

See also:

- Caterpillar Winners' Interview (<http://blog.kaggle.com/2015/09/22/caterpillar-winners-interview-1st-place-gilberto-josef-leustagos-mario/>)
- Caterpillar Winners' Interview: 3rd place (<http://blog.kaggle.com/2015/10/20/caterpillar-winners-interview-3rd-place-team-shift-workers/>)
- Liberty Mutual Property Inspection, Winner's Interview (<http://blog.kaggle.com/2015/09/28/liberty-mutual-property-inspection-winners-interview-qingchen-wang/>)
- Recruit Coupon Purchase Winner's Interview: 2nd place (<http://blog.kaggle.com/2015/10/21/recruit-coupon-purchase-winners-interview-2nd-place-halla-yang/>)
- Dato Truly Native? Winner's Interview: 2nd place (<http://blog.kaggle.com/2015/10/30/dato-winners-interview-2nd-place-mortehu/>)

```

Classifier: Gradient Boosting
Training time: 2409.8094s
Testing time: 0.4159s
Confusion matrix:
[[2214     1     3     5    10     8     9     3    24     1]
 [    1 2528    16    11     3     5     5     7     9     3]
 [    8     5 2165    34    16     5    12    22    37    10]
 [    1     9   27 2182     4    42     1    22    37    23]
 [    5     4    16     1 2088     5    12     5    10    65]
 [    9     6     7    41     8 1928    27     6    18    19]
 [   15     7     4     1    19    29 2181     1    14     0]
 [    6    16    27    15    22     6     0 2246     8    71]
 [    5    20    14    25    15    29     6     6 2057    38]
 [    6    10     8    24    49    15     1    54    17 2205]]
Accuracy: 0.9435

```

Naive Bayes ¶

Classifier: Naive Bayes
Training time: 0.3814s
Testing time: 0.8863s
Confusion matrix:
[[2094 4 11 10 6 7 56 3 69 18]
[4 2432 9 11 2 4 28 1 77 20]
[278 64 703 143 6 4 558 4 528 26]
[202 136 18 791 5 5 106 21 886 178]
[96 26 16 14 296 8 169 13 535 1038]
[327 63 15 39 14 87 100 5 1253 166]
[34 51 17 1 1 5 2109 0 52 1]
[19 21 3 23 20 2 7 737 123 1462]
[39 326 13 16 8 18 25 7 1482 281]
[15 26 8 2 14 2 1 41 40 2240]]
Accuracy: 0.5615

LDA ¶

Classifier: LDA
Training time: 20.6464s
Testing time: 0.0910s
Confusion matrix:
[[2131 2 10 14 12 47 20 4 36 2]
[0 2454 20 10 5 16 5 5 71 2]
[22 71 1873 77 51 8 82 20 101 9]
[5 32 56 1992 11 77 11 40 80 44]
[1 21 17 0 1983 12 12 2 21 142]
[19 18 11 112 18 1682 37 11 103 58]
[28 30 32 3 43 51 2046 0 37 1]
[16 57 25 20 70 8 0 1990 11 220]
[9 113 16 64 33 115 13 10 1781 61]
[15 10 6 35 133 14 0 122 22 2032]]
Accuracy: 0.8642

QDA ¶

Classifier: QDA
Training time: 23.0527s
Testing time: 6.2259s
Confusion matrix:
[[2212 3 12 14 1 4 20 5 6 1]
[66 2409 12 10 0 0 32 2 39 18]
[961 25 689 143 3 1 310 2 166 14]
[1231 48 29 606 3 13 66 10 232 110]
[810 22 25 27 250 4 143 17 345 568]
[909 15 13 33 1 214 140 4 666 74]
[83 18 14 1 1 2 2146 0 6 0]
[81 13 6 52 14 2 1 776 120 1352]
[487 181 18 20 6 17 58 3 1320 105]
[65 14 12 7 10 0 0 23 33 2225]]
Accuracy: 0.5561

MNIST Summary ¶

| Classifier | Accuracy | Training Time | Testing Time |
|----------------------|----------|---------------|--------------|
| MLP (500:200) | 97.98% | 79.5696s | 0.3480s |
| Dropout NN (500:200) | 97.80% | 118.2654s | 0.3918s |

| Classifier | Accuracy | Training Time | Testing Time |
|---|---------------|----------------|----------------|
| CNN (32 5×5 filters : 2×2 max pool : 64 5×5 filters : 2×2 max pool : 1024) | 97.69% | 391.8810s | 1.2035s |
| Adjusted SVM | 98.40% | 347.1539s | 234.5724s |
| Linear SVM | 94.16% | 168.6950s | 158.0101s |
| Random Forest (n_estimators=50, n_jobs=10) | 96.41% | 2.1359s | 26.0763s |
| Random Forest (n_estimators=10, max_features=1, max_depth=5) | 57.15% | 0.2077s | 22.2770s |
| k nearest neighbors (k=3) | 96.95% | 4.6439s | 1261.7815s |
| Decision Tree(max_depth=5) | 65.40% | 3.1346s | 0.0313s |
| Adaboost | 73.67% | 37.6443s | 1.5815s |
| Naive Bayes | 56.15% | 0.3814s | 0.8863s |
| LDA | 86.42% | 20.6464s | 0.0910s |
| QDA | 55.61% | 23.0527s | 6.2259s |
| Gradient Boosting | 94.35% | 2409.8094s | 0.4159s |
| Logistic Regression (C=1) | 91.47% | 272.1309s | 0.0531s |
| Logistic Regression (C=10000) | 91.23% | 1807.0624s | 0.0529s |

IRIS summary ¶

Just for fun, I tried the script from above with very minor adjustments to the IRIS flower dataset (https://en.wikipedia.org/wiki/Iris_flower_data_set):

| Classifier | Accuracy | Training Time | Testing Time |
|-------------------|---------------|---------------|----------------|
| AdaBoost | 92.00% | 0.1203s | 0.0101s |
| Decision Tree | 92.00% | 0.0005s | 0.0001s |
| Gradient Boosting | 92.00% | 0.2227s | 0.0007s |
| LDA | 96.00% | 0.0027s | 0.0002s |
| NN 20:5 | 90.00% | 1.6628s | 0.0046s |
| Naive Bayes | 90.00% | 0.0010s | 0.0004s |
| QDA | 94.00% | 0.0009s | 0.0003s |
| Random Forest | 90.00% | 0.2147s | 0.1395s |
| Random Forest 2 | 90.00% | 0.1481s | 0.1249s |
| SVM, adj. | 90.00% | 0.0010s | 0.0004s |
| SVM, linear | 88.00% | 0.0006s | 0.0002s |
| k nn | 92.00% | 0.0007s | 0.0009s |

| Classifier | Accuracy | Training Time | Testing Time |
|---|----------|---------------|----------------|
| Logistic Regression (C=1) | 88.00% | 0.0011s | 0.0001s |
| Logistic Regression (C=1000) | 92.00% | 0.0010s | 0.0002s |
| RBM 100 | 78.00% | 0.0233s | 0.0003s |
| RBM 100, n_iter=20 | 70.00% | 0.0427s | 0.0003s |
| RBM 200, n_iter=40, LR=0.01, Reg: C=1 | 88.00% | 0.2463s | 0.0005s |
| RBM 200, n_iter=40, LR=0.01, Reg: C=10000 | 90.00% | 0.2437s | 0.0005s |
| RBM 256 | 84.00% | 0.0424s | 0.0006s |
| RBM 512, n_iter=100 | 84.00% | 0.0723s | 0.0010s |

TL;DR ¶

Neural networks take their time to train and a feeling for the topology, but their classification results are nice and the testing time is good as well.

Random Forests and SVMs are also a model a type of model one should think of. However, the standard implementation is very slow compared to neural networks.

`sklearn lda.LDA` (<http://scikit-learn.org/0.16/modules/generated/sklearn lda.LDA.html>) might also be worth a try. The rest seems to be quite bad compared with those classifiers.

The code which generated the examples from above is here (<https://github.com/MartinThoma/algorithms/tree/master/ML/mnist/many-classifiers>).



Join the discussion...

LOG IN WITH

OR SIGN UP WITH DISQUS ?

Name

**Caterine** • a year ago

Hey, congrats for your blog! It is very good! =)

I have a question.. are all these methods implemented as pixel-based or feature-based?

^ | v • Reply • Share ›

**Martin Thoma** Mod → Caterine • a year ago

sklearn has general classifiers which work on feature vectors. In my case I chose each feature to be one color channel of one pixel; so each pixel gives 3 features.

^ | v • Reply • Share ›

**Caterine** → Martin Thoma • a year ago

Thanks Martin! Have you tried it with other feature representations such as SIFT or HOG descriptors?

^ | v • Reply • Share ›

**Martin Thoma** Mod → Caterine • a year ago

No.

^ | v • Reply • Share ›

[Subscribe](#) [Add Disqus to your site](#) [Add Disqus](#) [Disqus' Privacy Policy](#) [Privacy Policy](#) [Privacy Policy](#)

Published

Jan 19, 2016

by Martin Thoma (../author/martin-thoma/)

Category

Machine Learning (../categories.html#machine-learning-ref)

Tags

- Classification 6 (../tags.html#classification-ref)
- Machine Learning 70 (../tags.html#machine-learning-ref)
- Python 90 (../tags.html#python-ref)

Contact

[Twitter](https://twitter.com/themoosemind) (https://twitter.com/themoosemind)
 [Email](mailto:info@martin-thoma.de) (mailto:info@martin-thoma.de)
 [GitHub](https://github.com/MartinThoma) (https://github.com/MartinThoma)
 [Stack Overflow](http://stackoverflow.com/users/562769/martin-thoma) (http://stackoverflow.com/users/562769/martin-thoma)

- Martin Thoma - A blog about Code, the Web and Cyberculture
- Datenschutzerklärung (http://www.martin-thoma.de/privacy.htm)
- Impressum (http://www.martin-thoma.de/impressum.htm)
- Powered by Pelican (http://getpelican.com/). Theme: Elegant (http://oncrashreboot.com/pelican-elegant) by Talha Mansoor (http://oncrashreboot.com)